

Code challenge easynvest

Description

This is the complete document of the challenge proposed by Easynvest. Disclosure of company name and publishing of the results were explicitly authorized by their recruiting team.

Easynvest is a fintech company, more specifically a digital broker-dealer which helps thousands of clients to invest their money easily and quickly. They are known for their online platform and strong digital presence.

The complete description of the challenge may be found in the file `challenge_description/challenge_description.pdf` (in Portuguese).

Introduction

The data set

The data set received was in the form of an Excel spreadsheet with two tabs. The first tab contained 4973 entries ($N=4973$), one unique ID and 10 characteristics (11 columns total).

The second tab has entries which are not described elsewhere. The lack of a formal description casts unnecessary uncertainty into the data at hand. Lack of proper definition is a discouraged practice in data creation (e.g.: absence of research method and methodology). Data definition must not be open for interpretation.

A remarkable fact of this data set is that it does not contain any null values. Such high quality data sets are rare to find and may indicate that its source is very thoughtful of its data management.

A final remark is that the characteristics' names (column names) should not be considered self explanatory. A codebook is often use to describe published data. To illustrate the critique above consider the variable 'VALOR_01' (value_01; there are 4 of these variables). To what value does it refer to? Is it the amount already invested in the investment platform? Is it the income enumerated by different sources of income? Is it profit? If it is income, is it yearly or monthly? Another illustration is the 'GEO_REFERENCIA' (georeference) variable. It has values ranging from 10 to 999 but it is not explained elsewhere. Usual geolocation information are comprised of x and y coordinates or other better known formats. Also, it cannot refer to Brazilian municipalities because there exists ~5500 of them.

Consequently this variable has been neglected in the present analysis.

As one can see, this seemingly unimportant differences may yield different interpretations later on the data analysis and render some conclusions useless or even worse: wrong.

Approach

As stated in the challenge description my work should:

1. **Group users finding well defined groups with common characteristics.**
 - In order to do that I have clustered the data set using the *K-Means* clustering algorithm.
2. **Justify the chosen clustering algorithm.**
 - This algorithm is one of the most *commonly used* algorithm in Data Sciences. As such one can easily find support, implementations, discussions and suggestions on various references. Such vast amount of information is not something to be neglected.
The algorithm also *allows the specification of the number of clusters* to be found. This is seen as drawback according to some sources. Yet I think that it can be overcome with successively running the algorithm with a different cluster number. Specifying the number of clusters also impedes the algorithm to come up with a number of clusters which may be uninterpretable (too few, e. g. 2 or too many 10+).
The algorithm tends to yield *clusters with similar size*. This may be a desired characteristic in a business setting for example, where investment of resources (time and capital) may be applied to a cluster of clients. In such cases one does not want to invest those in a cluster just to find out that it aggregates to just a few individuals of their clientele.
3. **Present metrics of performance for the chosen algorithm.**
 - In this case the *silhouette analysis* was performed to assess the effectiveness of the clustering algorithm.
Also the *intra-group and inter-group standard deviation* and means were taken in consideration to interpret the results of this clustering algorithm.
4. **Discuss the metrics of performance to assess the clusters.**
 - See discussion of the clustering for a detailed assessment of the clustering algorithm.
5. **Explain the results.**
 - See the results and summary sections for a precise answer to this question.

Results

Preprocessing

Variable scaling

The received data needed preprocessing before applying the clustering method. That is because the K-Means clustering method is sensitive to variable scaling (more precisely to variance). Without scaling, variables tend to have variances of different orders of magnitude (standard deviation for the data set before preprocessing):

variable	std
valor_01	6098.823
valor_02	89180.835
valor_03	37645.943
valor_04	23246.037
age	10.792
estado_civil_solteiro	0.500
estado_civil_casado	0.486
estado_civil_outro	0.297
genero_m	0.416
genero_f	0.416
perfil_a	0.458
perfil_b	0.416
perfil_c	0.216
perfil_d	0.161

Standard deviation for the data set after preprocessing (abbreviated):

variable	std
valor_01	1.0
valor_02	1.0
(...)	1.0
perfil_c	1.0
perfil_d	1.0

That means that without scaling the four variables of 'valor' would dominate the clustering sensitivity, rendering the presence of the other variables useless.

Nominal variables processing

Some presented variables are categorical and do not meaningfully present any interpretation from a numerical standpoint. For example, height may be compared so that a person who is 170 cm high is higher than someone who is 165 cm.

There is no parallel to variables which represent ‘non rankable’ variables such as gender and ethnicity. Assigning a value of 1 for male, 0 for female and 2 for non identified gender does not mean that in this scenario that male > female.

In order to overcome this problem categorical variables with N categories are transformed to new binary characteristics (then scaled as commented above). To illustrate suppose that we begin only with `col1` and `col_a`, `col_b` and `col_c` are generated from them:

<code>col1</code>	<code>col_a</code>	<code>col_b</code>	<code>col_c</code>
a	1	0	0
b	0	1	0
c	0	0	1
a	1	0	0

This allow them to be included in the K-Means clustering algorithm.

Clustering

Choice of the number of clusters

I have chosen the numbers of clusters to be six. See the discussion below for details.

Before diving in the details of my choice, one cannot overstress the importance of the choice of the number of clusters. This is arguably the most tricky decision in this challenge as it deals with a great mix of technical as well as non-technical details.

Silhouette analysis (technical analysis)

Silhouette analysis is a technique used to compare how well your data is sorted into clusters. It can be calculated to all data points and then averaged to provide a summary statistic. It ranges from -1 to 1:

- Values near to -1: the data point was incorrectly clustered and should belong to a different cluster
- Values near to zero: the point lies between two clusters and lack a sharp ‘belonging attribute’ (it could thus belong to both clusters)

- Values near to one: the data point was correctly classified and lies near to other data points in the same cluster. Its cluster is adequately away from other clusters

From a pure technical standpoint choosing the number of clusters such that the average value for silhouette is maximum is the best option. On the other hand, working with such a large number of clusters may hinder the interpretability of the results as clusters probably would not have a sharp distinction between them (consider that our data set has 10 dimensions originally). Probably the communication of such results for a multidisciplinary team of mixed background would be noisy as well.

Real world analysis (non-technical analysis)

I have chosen the number of cluster to be 6 for a couple of different reasons. First of all, analyzing the average value for silhouette we can see that the average value for silhouette reaches a maximum at around 18 clusters. Thus we naively could choose the number of clusters to be 18.

However in the context of the **interpretability and communications** of the results one would limit the number of clusters to a maximum of ~10.

Back to the average silhouettes, we can see that it is an increasing function between 2 and 6 clusters, almost doubling its value in this interval. This means that the samples are on average better defined in their own cluster and far away from other clusters. Another fact that indicates that 6 is a good number for clusters is that in this case just a few data points show a silhouette smaller than zero. In other words, just a few data points are incorrectly labeled in their cluster (those data points are unfrequent and are concentrated on cluster 2) (see below). Using the same argumentation the cluster that is best defined is cluster 1 because of the high incidence of data points at near 0.75 silhouette value.

See images for silhouette for all images.

Cluster interpretation

For cluster interpretation two resources are available:

1. Tables output to stdout during program execution.
2. Plots.

Please notice that all tables, plots and results are exhaustively detailed below. They are repeated here for convenience.

From a general standpoint clusters should have low intra-cluster variance and high inter-cluster variance for each variable.




Figure 1:

Cluster 0

Distinctive features:

1. Has all individuals with 'other' marital status (that is, it is neither married nor single).




Figure 2:

1. Has the highest age mean of all groups even though there is a high dispersion both intra and inter cluster for this variable.

Cluster 1

Distinctive features:

1. Has the most concentration of single persons ('solteiro').
1. It is solely composed of male individuals (absence of 'genenro_f'). This also happens to cluster 2 and cluster 3.
2. Has the most concentration of profile D ('perfil_d'). Also contains a lot of profile A individuals.
3. Has the most concentration young people.
4. Has the highest average value of silhouette (see above).
5. It is the cluster which aggregates most individuals (~1400).




Figure 3:

Cluster 2

Distinctive features:

1. Has the highest averages for 'valor_02', 'valor_03' and 'valor_04'. In respect to these 3 variables all the other groups have lower averages.
1. Includes almost solely profile B people.
2. Contains almost solely males.

Cluster 3

Distinctive features:

1. It is the group with the highest proportion of married individuals ('estado_civil_casado'). And it is solely comprised of married individuals.
1. The cluster is entirely comprised of male individuals.
2. The cluster contains only individuals from profile A and profile D.

Cluster 4

Distinctive features:




Figure 4:




Figure 5:

1. Comprised solely of female subjects:




Figure 6:

1. Contains a fair distribution of different profiles ('perfil'), marital status, and values ('valor')

Cluster 5

Distinctive features:

1. Comprised solely of profile C and it contains this group entirely:

(the absence of the errorbar indicates that there is only one value for this variable in for this cluster).

1. It is the smallest of all clusters: 245 individuals.

Summary

The clusterization was conducted properly and yielded significant results. This is evidenced by:

- A satisfactory value of silhouette indicated good clustering (0.433)




Figure 7:

- A reasonable amount of clusters, facilitating the communication and interpretation of the results (one of the strenghts of the algorithm)
- Some very sharp separations, some of which are coupled and yield easily interpretable results:
 - Cluster 0:
 - * it contains all individuals with ‘other’ marital status
 - * has the highest age mean
 - * (notice that the two variables are correlated)
 - Cluster 1:
 - * all individuals are single males
 - * lowest age mean of all clusters
 - * it is the largest cluster
 - Cluster 2:
 - * concentrates individuals with high value variables (‘valor’) from 2 to 4
 - Cluster 3:
 - * group entirely comprised of male individuals
 - * group entirely comprised of married individuals
 - Cluster 4:
 - * group entirely composed of female subjects
 - Cluster 5:

* it contains all individuals with profile C ('perfil_c')

Additional information & Reproducibility

Reproducibility

Reproducibility is going to be assessed in this task. In order to comply with it the software versions needed to replicate the experiment are specified below.

Also non deterministic part of the algorithms are fixed using a defined random seed at `code/control.py` and invoked properly during code execution.

Finally the code is hosted on github to allow any team to replicate and judge the results themselves.

Tools

- Vim

```
vim --version
VIM - Vi IMproved 8.0 (2016 Sep 12, compiled Apr  4 2017 13:41:19)
Included patches: 1-542
Modified by <cygwin@cygwin.com>
Compiled by <cygwin@cygwin.com>
Huge version without GUI.
```

- python

```
python3 --version
Python 3.6.1
```

- python modules:

```
data-utilities==1.2.6
matplotlib==2.0.0
numpy==1.12.1
pandas==0.19.2
scikit-learn==0.18.1
scipy==0.19.0
seaborn==0.7.1
```

- pandoc

```
pandoc --version
pandoc.exe 1.19.2.1
Compiled with pandoc-types 1.17.0.4, texmath 0.9, skylighting 0.1.1.4
Default user data directory: C:\Users\e061568\AppData\Roaming\pandoc
```

Copyright (C) 2006-2016 John MacFarlane
Web: <http://pandoc.org>
This is free software; see the source for copying conditions.
There is no warranty, not even for merchantability or fitness
for a particular purpose.

Other remarks

Here are comments which would not fit elsewhere in the discussion of the document. Despite there are discussions on some of these topics I thought they would not fit well in the flow of the assignment.

- As suggested in the introduction, a more precise definition of the data set could improve the conclusions that could be drawn from it. When presenting a data set for someone who is not acquainted with how it was generated extreme care should be taken in order to communicate the variables, their origin and their precise meaning carefully.
- K-Means clusterization algorithm:
 - Randomization of the initial cluster points may yield very different clusters for the given data. This is considered a weakness. To overcome this the algorithm may use the *k-means++* initial seeding which improve the initial assignment of the algorithm.
- Approximate running time for the `main.py` is 7 minutes on commodity hardware:

```
time python3 main.py
7:06.02
```

It is dominated by the plotting of different k-means silhouettes. The rest of the code runs in less than one minute (full running time is kept for reproducibility).

A possible speedup would be to run the KMeans object with the keyword argument `n_jobs=-1` but this causes cygwin (the platform in which this code is being developed) to halt. Possibly in normal windows/linux environments this would not happen.

Next steps

- No more next steps are needed. The code is in its final version as well as the output file.

All output from python code

All the images

Silhouette




Figure 8: file silhouette.png

Clusters

Notice that the error bars represented here are ± 1 standard deviation.

Code output (stdout)

```
-----  
The effect of preprocessing on standard deviation  -----  
-----
```

```
    before:  
valor_01          6098.823  
valor_02          89180.835  
valor_03          37645.943  
valor_04          23246.037
```




Figure 9: file silhouette_120.png



Figure 10: file silhouette_distribution_for_n=2_clusters.png




Figure 11: file silhouette_distribution_for_n=3_clusters.png




Figure 12: file silhouette_distribution_for_n=4_clusters.png




Figure 13: file silhouette_distribution_for_n=5_clusters.png




Figure 14: file silhouette_distribution_for_n=6_clusters.png




Figure 15: file silhouette_distribution_for_n=7_clusters.png




Figure 16: file cluster_analysis_cluster_mean_of_variables_age.png




Figure 17: file cluster_analysis_cluster_mean_of_variables_data_nascimento_int.png




Figure 18: file cluster_analysis_cluster_mean_of_variables_estado_civil_casado.png




Figure 19: file cluster_analysis_cluster_mean_of_variables_estado_civil_outro.png




Figure 20: file cluster_analysis_cluster_mean_of_variables_estado_civil_solteiro.png




Figure 21: file cluster_analysis_cluster_mean_of_variables_genero_f.png




Figure 22: file cluster_analysis_cluster_mean_of_variables_genero_m.png




Figure 23: file cluster_analysis_cluster_mean_of_variables_geo_referencia.png




Figure 24: file cluster_analysis_cluster_mean_of_variables_id.png




Figure 25: file cluster_analysis_cluster_mean_of_variables_perfil_a.png




Figure 26: file cluster_analysis_cluster_mean_of_variables_perfil_b.png




Figure 27: file cluster_analysis_cluster_mean_of_variables_perfil_c.png




Figure 28: file cluster_analysis_cluster_mean_of_variables_perfil_d.png




Figure 29: file cluster_analysis_cluster_mean_of_variables_valor_01.png




Figure 30: file cluster_analysis_cluster_mean_of_variables_valor_02.png




Figure 31: file cluster_analysis_cluster_mean_of_variables_valor_03.png




Figure 32: file cluster_analysis_cluster_mean_of_variables_valor_04.png

```

age                                10.792
estado_civil_solteiro              0.500
estado_civil_casado                0.486
estado_civil_outro                 0.297
genero_m                           0.416
genero_f                           0.416
perfil_a                           0.458
perfil_b                           0.416
perfil_c                           0.216
perfil_d                           0.161
dtype: float64
after:
valor_01                           1.0
valor_02                           1.0
valor_03                           1.0
valor_04                           1.0
age                                1.0
estado_civil_solteiro              1.0
estado_civil_casado                1.0
estado_civil_outro                 1.0
genero_m                           1.0
genero_f                           1.0

```

```

perfil_a          1.0
perfil_b          1.0
perfil_c          1.0
perfil_d          1.0
dtype: float64

```

```

-----
Features for different number of clusters  -----
-----

```

```

Silhouette average for 2 clusters 0.256

```

```

Number of individuals per cluster for 2 clusters -----

```

```

(2392,)

```

```

(2580,)

```

```

Silhouette average for 3 clusters 0.310

```

```

Number of individuals per cluster for 3 clusters -----

```

```

(1887,)

```

```

(1979,)

```

```

(1106,)

```

```

Silhouette average for 4 clusters 0.349

```

```

Number of individuals per cluster for 4 clusters -----

```

```

(1103,)

```

```

(1492,)

```

```

(1405,)

```

```

(972,)

```

```

Silhouette average for 5 clusters 0.374

```

```

Number of individuals per cluster for 5 clusters -----

```

```

(1431,)

```

```

(972,)

```

```

(1072,)

```

```

(133,)

```

```

(1364,)

```

```

Silhouette average for 6 clusters 0.433

```

```

Number of individuals per cluster for 6 clusters -----

```

```

(451,)

```

```

(1407,)

```

```

(884,)

```

```

(1062,)

```

```

(923,)

```

```

(245,)

```

```

Silhouette average for 7 clusters 0.426

```

```

Number of individuals per cluster for 7 clusters -----

```

```

(23,)

```

```

(1547,)

```

```

(1017,)

```

```

(952,)

```

```

(133,)

```

(244,)
(1056,)
Silhouette average for 8 clusters 0.468
Silhouette average for 9 clusters 0.485

Cluster standard deviation of variables -----

	id	geo_referencia	valor_01	valor_02	valor_03	\
cluster						
0	1432.085	318.310	18855.499	50776.256	17566.300	
1	1444.033	301.413	1273.152	21754.206	3742.253	
2	1429.445	300.858	3528.072	178904.471	83183.573	
3	1440.205	298.530	2208.455	53021.604	9307.479	
4	1431.331	294.305	1391.051	53714.303	5459.086	
5	1417.395	291.092	3051.785	104054.907	44757.829	
all	1435.437	300.712	6098.823	89180.835	37645.943	

	valor_04	data_nascimento_int	age	estado_civil_solteiro	\
cluster					
0	27998.787	3.714e+17	11.769		0.000
1	8408.652	1.972e+17	6.248		0.000
2	33028.168	3.354e+17	10.628		0.498
3	26499.920	3.286e+17	10.414		0.000
4	21677.157	3.467e+17	10.985		0.484
5	16546.183	3.947e+17	12.508		0.499
all	23246.037	3.406e+17	10.792		0.500

	estado_civil_casado	estado_civil_outro	genero_m	genero_f	\
cluster					
0	0.000	0.000	0.450	0.450	
1	0.000	0.000	0.000	0.000	
2	0.497	0.048	0.058	0.058	
3	0.000	0.000	0.000	0.000	
4	0.484	0.000	0.000	0.000	
5	0.494	0.338	0.415	0.415	
all	0.486	0.297	0.416	0.416	

	perfil_a	perfil_b	perfil_c	perfil_d
cluster				
0	0.433	0.422	0.000	0.132
1	0.204	0.000	0.000	0.204
2	0.034	0.034	0.000	0.000
3	0.186	0.000	0.000	0.186
4	0.363	0.334	0.000	0.166
5	0.000	0.000	0.000	0.000

all	0.458	0.416	0.216	0.161
-----	-------	-------	-------	-------

Cluster mean of variables -----

	id	geo_referencia	valor_01	valor_02	valor_03	valor_04	\
cluster							
0	2452.268	360.517	3176.098	16914.336	3676.908	6541.581	
1	2524.927	336.119	1241.121	5949.802	472.001	2373.197	
2	2457.329	339.374	2737.195	49328.855	15918.938	8769.796	
3	2469.815	349.920	2329.157	15720.792	1585.815	5992.588	
4	2497.011	320.521	1448.414	10067.777	631.819	3835.935	
5	2466.812	292.376	2645.085	28873.244	10001.609	4562.969	
all	2486.500	336.808	2022.698	18638.059	4246.265	5041.123	

	data_nascimento_int	age	estado_civil_solteiro	\
cluster				
0	1.856e+17	41.120	0.000	
1	5.707e+17	28.918	1.000	
2	3.262e+17	36.665	0.552	
3	2.436e+17	39.281	0.000	
4	4.143e+17	33.871	0.626	
5	2.291e+17	39.741	0.453	
all	3.766e+17	35.068	0.520	

	estado_civil_casado	estado_civil_outro	genero_m	genero_f	\
cluster					
0	0.000	1.000	0.718	0.282	
1	0.000	0.000	1.000	0.000	
2	0.446	0.002	0.997	0.003	
3	1.000	0.000	1.000	0.000	
4	0.374	0.000	0.000	1.000	
5	0.416	0.131	0.780	0.220	
all	0.383	0.098	0.777	0.223	

	perfil_a	perfil_b	perfil_c	perfil_d
cluster				
0	0.752	0.231	0.000	0.018
1	0.957	0.000	0.000	0.043
2	0.001	0.999	0.000	0.000
3	0.964	0.000	0.000	0.036
4	0.844	0.128	0.000	0.028
5	0.000	0.000	1.000	0.000
all	0.702	0.222	0.049	0.027


```

Normalized cluster standard deviation of variables  -----
-----
      id  geo_referencia  valor_01  valor_02  valor_03  valor_04  \
cluster
0      0.992           1.000      1.000      0.284      0.211      0.848
1      1.000           0.947      0.068      0.122      0.045      0.255
2      0.990           0.945      0.187      1.000      1.000      1.000
3      0.997           0.938      0.117      0.296      0.112      0.802
4      0.991           0.925      0.074      0.300      0.066      0.656
5      0.982           0.914      0.162      0.582      0.538      0.501
all    0.994           0.945      0.323      0.498      0.453      0.704

```

```

      data_nascimento_int  age  estado_civil_solteiro  \
cluster
0              0.941  0.941              0.000
1              0.499  0.499              0.000
2              0.850  0.850              0.996
3              0.833  0.833              0.000
4              0.878  0.878              0.969
5              1.000  1.000              0.998
all            0.863  0.863              1.000

```

```

      estado_civil_casado  estado_civil_outro  genero_m  genero_f  \
cluster
0              0.000              0.000      1.000      1.000
1              0.000              0.000      0.000      0.000
2              1.000              0.141      0.129      0.129
3              0.000              0.000      0.000      0.000
4              0.973              0.000      0.000      0.000
5              0.993              1.000      0.922      0.922
all            0.977              0.879      0.924      0.924

```

```

      perfil_a  perfil_b  perfil_c  perfil_d
cluster
0      0.945      1.000      0.0      0.649
1      0.445      0.000      0.0      1.000
2      0.074      0.080      0.0      0.000
3      0.406      0.000      0.0      0.912
4      0.794      0.792      0.0      0.813
5      0.000      0.000      0.0      0.000
all    1.000      0.986      1.0      0.792

```

```

-----
Normalized cluster mean  -----
-----
      id  geo_referencia  valor_01  valor_02  valor_03  valor_04  \

```

cluster						
0	0.971	1.000	1.000	0.343	0.231	0.746
1	1.000	0.932	0.391	0.121	0.030	0.271
2	0.973	0.941	0.862	1.000	1.000	1.000
3	0.978	0.971	0.733	0.319	0.100	0.683
4	0.989	0.889	0.456	0.204	0.040	0.437
5	0.977	0.811	0.833	0.585	0.628	0.520

	data_nascimento_int	age	estado_civil_solteiro	\
cluster				
0	0.325	1.000		0.000
1	1.000	0.703		1.000
2	0.572	0.892		0.552
3	0.427	0.955		0.000
4	0.726	0.824		0.626
5	0.401	0.966		0.453

	estado_civil_casado	estado_civil_outro	genero_m	genero_f	\
cluster					
0	0.000	1.000	0.718	0.282	
1	0.000	0.000	1.000	0.000	
2	0.446	0.002	0.997	0.003	
3	1.000	0.000	1.000	0.000	
4	0.374	0.000	0.000	1.000	
5	0.416	0.131	0.780	0.220	

	perfil_a	perfil_b	perfil_c	perfil_d
cluster				
0	0.780	0.231	0.0	0.409
1	0.992	0.000	0.0	1.000
2	0.001	1.000	0.0	0.000
3	1.000	0.000	0.0	0.825
4	0.875	0.128	0.0	0.650
5	0.000	0.000	1.0	0.000

Cluster sum of variables -----

	id	geo_referencia	valor_01	valor_02	valor_03	\
cluster						
0	1.106e+06	162593.0	1.432e+06	7.628e+06	1.658e+06	
1	3.553e+06	472919.0	1.746e+06	8.371e+06	6.641e+05	
2	2.172e+06	300007.0	2.420e+06	4.361e+07	1.407e+07	
3	2.623e+06	371615.0	2.474e+06	1.670e+07	1.684e+06	
4	2.305e+06	295841.0	1.337e+06	9.293e+06	5.832e+05	
5	6.044e+05	71632.0	6.480e+05	7.074e+06	2.450e+06	

	valor_04	data_nascimento_int	age	estado_civil_solteiro	\
cluster					
0	2.950e+06	8.370e+19	18544.975		0.0
1	3.339e+06	8.029e+20	40687.039		1407.0
2	7.752e+06	2.883e+20	32411.595		488.0
3	6.364e+06	2.587e+20	41715.986		0.0
4	3.541e+06	3.824e+20	31263.225		578.0
5	1.118e+06	5.613e+19	9736.619		111.0

	estado_civil_casado	estado_civil_outro	genero_m	genero_f	\
cluster					
0	0.0	451.0	324.0	127.0	
1	0.0	0.0	1407.0	0.0	
2	394.0	2.0	881.0	3.0	
3	1062.0	0.0	1062.0	0.0	
4	345.0	0.0	0.0	923.0	
5	102.0	32.0	191.0	54.0	

	perfil_a	perfil_b	perfil_c	perfil_d
cluster				
0	339.0	104.0	0.0	8.0
1	1346.0	0.0	0.0	61.0
2	1.0	883.0	0.0	0.0
3	1024.0	0.0	0.0	38.0
4	779.0	118.0	0.0	26.0
5	0.0	0.0	245.0	0.0

Bibliography

K-Means algorithm

1. https://en.wikipedia.org/wiki/K-means_clustering
2. <http://scikit-learn.org/stable/modules/generated/sklearn.cluster.KMeans.html#sklearn.cluster.KMeans>

Silhouette analysis

1. <http://scikit-learn.org/stable/modules/clustering.html#silhouette-coefficient>
2. <http://www.sciencedirect.com/science/article/pii/0377042787901257>

Preprocessing

1. <http://scikit-learn.org/stable/modules/preprocessing.html>
2. <https://stats.stackexchange.com/questions/21222/are-mean-normalization-and-feature-scaling-needed-for-k-means-clustering>